

```
In [1]: import pandas as pd

# Load the CSV file
file_path = 'dataset_traffic_accident_prediction1.csv' # Use this exact name
df = pd.read_csv(file_path)

# Show the shape of the dataset
print("Shape of the dataset:", df.shape)

# Display column names
print("Columns:\n", df.columns.tolist())

# Preview the first 5 rows
df.head()
```

Shape of the dataset: (840, 14)

Columns:

['Weather', 'Road\_Type', 'Time\_of\_Day', 'Traffic\_Density', 'Speed\_Limit', 'Number\_of\_Vehicles', 'Driver\_Alcohol', 'Accident\_Severity', 'Road\_Condition', 'Vehicle\_Type', 'Driver\_Age', 'Driver\_Experience', 'Road\_Light\_Condition', 'Accident']

```
Out[1]:
```

	Weather	Road_Type	Time_of_Day	Traffic_Density	Speed_Limit	Number_of_Vehicles	Driver_Alcohol
0	Rainy	City Road	Morning	1.0	100.0	5.0	
1	Clear	Rural Road	Night	NaN	120.0	3.0	
2	Rainy	Highway	Evening	1.0	60.0	4.0	
3	Clear	City Road	Afternoon	2.0	60.0	3.0	
4	Rainy	Highway	Morning	1.0	195.0	11.0	

```
In [5]: # Check for missing values
print("Missing values per column:\n", df.isnull().sum())

# Drop columns with too many missing values (threshold optional)
df = df.dropna(thresh=0.5*len(df), axis=1)

# Replace this:
df.fillna(method='ffill', inplace=True)

# ✓ With this:
df.ffill(inplace=True)

# Verify
print("Remaining missing values:", df.isnull().sum().sum())
```

Missing values per column:

```

Weather      0
Road_Type    0
Time_of_Day  0
Traffic_Density 0
Speed_Limit  0
Number_of_Vehicles 0
Driver_Alcohol 0
Accident_Severity 1
Road_Condition 0
Vehicle_Type 0
Driver_Age    0
Driver_Experience 0
Road_Light_Condition 0
Accident      0
dtype: int64

```

Remaining missing values: 1

C:\Users\SHANGAR\AppData\Local\Temp\ipykernel\_13180\550429416.py:8: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
df.fillna(method='ffill', inplace=True)
```

```

In [4]: # Check for missing values
print("Missing values per column:\n", df.isnull().sum())

# Drop columns with too many missing values (threshold optional)
df = df.dropna(thresh=0.5*len(df), axis=1)

# Fill missing values
df.fillna(method='ffill', inplace=True) # Forward fill as default
# OR: df.fillna(df.mean(numeric_only=True), inplace=True) # Numeric only

# Verify
print("Remaining missing values:", df.isnull().sum().sum())

```

Missing values per column:

```

Weather      0
Road_Type    0
Time_of_Day  0
Traffic_Density 0
Speed_Limit  0
Number_of_Vehicles 0
Driver_Alcohol 0
Accident_Severity 1
Road_Condition 0
Vehicle_Type 0
Driver_Age    0
Driver_Experience 0
Road_Light_Condition 0
Accident      0
dtype: int64

```

Remaining missing values: 1

C:\Users\SHANGAR\AppData\Local\Temp\ipykernel\_13180\124349902.py:8: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
df.fillna(method='ffill', inplace=True) # Forward fill as default
```

```
In [6]: import pandas as pd

# Load dataset
df = pd.read_csv('dataset_traffic_accident_prediction1.csv')

# Step 1: Check missing values
print("🔍 Missing values before handling:\n")
print(df.isnull().sum())

# Step 2: Drop columns with too many missing values (optional)
# Drops columns where more than 50% of values are missing
df = df.dropna(thresh=0.5 * len(df), axis=1)

# Step 3: Forward fill remaining missing values (recommended method)
df.ffill(inplace=True)

# OPTIONAL: You could also backward fill if needed
# df.bfill(inplace=True)

# Step 4: Check again
print("\n✅ Missing values after handling:\n")
print(df.isnull().sum().sum()) # Should be 0 if all handled
```

🔍 Missing values before handling:

```
Weather          42
Road_Type        42
Time_of_Day      42
Traffic_Density  42
Speed_Limit      42
Number_of_Vehicles 42
Driver_Alcohol   42
Accident_Severity 42
Road_Condition   42
Vehicle_Type     42
Driver_Age       42
Driver_Experience 42
Road_Light_Condition 42
Accident         42
dtype: int64
```

✅ Missing values after handling:

1

```
In [11]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the cleaned dataset
df = pd.read_csv('dataset_traffic_accident_prediction1.csv')


# Optional: Fill or drop missing values before plotting if not already done
df.ffill(inplace=True)

# Set visual style
```

```
sns.set(style='whitegrid')

# Compute correlation matrix
correlation_matrix = df[numeric_cols].corr()

# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=
plt.title('Correlation Heatmap of Numeric Features')
plt.tight_layout()
plt.show()
```



In [ ]:

```
In [17]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load and prepare the dataset
df = pd.read_csv('dataset_traffic_accident_prediction1.csv')
df.ffill(inplace=True) # Handle missing values

# Example: encode target variable (adjust column name if needed)
if df['Accident_Severity'].dtype == 'object':
    le = LabelEncoder()
```

```

df['Accident_Severity'] = le.fit_transform(df['Accident_Severity'])

# Optional: One-hot encode other categorical features
df = pd.get_dummies(df, drop_first=True)

# Define features (X) and target (y)
X = df.drop('Accident_Severity', axis=1)
y = df['Accident_Severity']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 3. XGBoost (Advanced)
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss', random_state=42)
xgb_model.fit(X_train, y_train)
xgb_preds = xgb_model.predict(X_test)
print("XGBoost Accuracy:", accuracy_score(y_test, xgb_preds))
print(classification_report(y_test, xgb_preds))

```

C:\Users\SHANGAR\anaconda3\Lib\site-packages\xgboost\training.py:183: UserWarning: [12:05:08] WARNING: C:\actions-runner\\_work\xgboost\xgboost\src\learner.cc:738: Parameters: { "use\_label\_encoder" } are not used.

```

bst.update(dtrain, iteration=i, fobj=obj)
XGBoost Accuracy: 0.5059523809523809

```

	precision	recall	f1-score	support
0	0.25	0.06	0.10	17
1	0.58	0.74	0.65	99
2	0.28	0.21	0.24	52
accuracy			0.51	168
macro avg	0.37	0.34	0.33	168
weighted avg	0.46	0.51	0.47	168

```

In [19]: # Assumes models and predictions are already made from previous steps
# y_test: true labels
# log_preds, rf_preds, xgb_preds: predictions from different models
from sklearn.metrics import classification_report


from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Confusion matrix
cm = confusion_matrix(y_test, rf_preds)

# Plot heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')

```

```
plt.tight_layout()
plt.show()
```



```
In [20]: from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.preprocessing import label_binarize

# If multiclass, binarize the output
y_test_bin = label_binarize(y_test, classes=[0, 1, 2]) # Adjust based on your target
rf_preds_proba = rf_model.predict_proba(X_test)

# Plot ROC curve for each class
plt.figure(figsize=(8, 6))
for i in range(y_test_bin.shape[1]):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], rf_preds_proba[:, i])
    auc = roc_auc_score(y_test_bin[:, i], rf_preds_proba[:, i])
    plt.plot(fpr, tpr, label=f'Class {i} (AUC = {auc:.2f})')

plt.plot([0, 1], [0, 1], 'k--') # Random Line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Random Forest')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

